

Project-2 Continuous Control

23RD MAR 2020

AMRIT PAUL



Problem Statement

The goal of the project is to demonstrate the abilities of a model-free reinforcement learning algorithm, particularly Deep Deterministic Policy Gradients (DDPG) Algorithm, which consists of two neural networks namely Actor network and Critic network. The project uses Unity environment, a game development framework and Pytorch, a deep learning framework. The algorithm was trained on 20 agents, with each agent consisting of 33 states and 4 actions.

Description

Deep Deterministic Policy Gradient (DDPG) is an algorithm which concurrently learns a Q-function and a policy. It uses off-policy data and the Bellman equation to learn the Q-function and uses the Q-function to learn the policy.

The approach is closely connected to Q-learning and is motivated the same way: if you know the optimal action-value function $Q^*(s,a)$, then in any given state, the optimal action $a^*(s)$ can be found by solving :-

$$a^*(s) = \arg \max_a Q^*(s, a).$$

DDPG interleaves learning an approximator to $Q^*(s,a)$ with learning an approximator to $a^*(s)$ and it does so in a way which is specifically adapter for environments with continuous action spaces. The actor produces a deterministic policy instead of the usual stochastic policy and the critic evaluates the deterministic policy. The critic is updated using the TD-error and the actor is trained using the deterministic policy gradient algorithm.

$$\begin{aligned}\nabla_{\theta^\mu} J &\approx \mathbb{E}_{s_t \sim \rho^\beta} [\nabla_{\theta^\mu} Q(s, a | \theta^Q) |_{s=s_t, a=\mu(s_t | \theta^\mu)}] \\ &= \mathbb{E}_{s_t \sim \rho^\beta} [\nabla_a Q(s, a | \theta^Q) |_{s=s_t, a=\mu(s_t)} \nabla_{\theta^\mu} \mu(s | \theta^\mu) |_{s=s_t}]\end{aligned}$$

Pseudocode

Algorithm 1 Deep Deterministic Policy Gradient

```
1: Input: initial policy parameters  $\theta$ , Q-function parameters  $\phi$ , empty replay buffer  $\mathcal{D}$ 
2: Set target parameters equal to main parameters  $\theta_{\text{targ}} \leftarrow \theta$ ,  $\phi_{\text{targ}} \leftarrow \phi$ 
3: repeat
4:   Observe state  $s$  and select action  $a = \text{clip}(\mu_{\theta}(s) + \epsilon, a_{\text{Low}}, a_{\text{High}})$ , where  $\epsilon \sim \mathcal{N}$ 
5:   Execute  $a$  in the environment
6:   Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
7:   Store  $(s, a, r, s', d)$  in replay buffer  $\mathcal{D}$ 
8:   If  $s'$  is terminal, reset environment state.
9:   if it's time to update then
10:    for however many updates do
11:      Randomly sample a batch of transitions,  $B = \{(s, a, r, s', d)\}$  from  $\mathcal{D}$ 
12:      Compute targets
```

$$y(r, s', d) = r + \gamma(1 - d)Q_{\phi_{\text{targ}}}(s', \mu_{\theta_{\text{targ}}}(s'))$$

```
13:   Update Q-function by one step of gradient descent using
```

$$\nabla_{\phi} \frac{1}{|B|} \sum_{(s, a, r, s', d) \in B} (Q_{\phi}(s, a) - y(r, s', d))^2$$

```
14:   Update policy by one step of gradient ascent using
```

$$\nabla_{\theta} \frac{1}{|B|} \sum_{s \in B} Q_{\phi}(s, \mu_{\theta}(s))$$


```
15:   Update target networks with
```

$$\begin{aligned}\phi_{\text{targ}} &\leftarrow \rho \phi_{\text{targ}} + (1 - \rho) \phi \\ \theta_{\text{targ}} &\leftarrow \rho \theta_{\text{targ}} + (1 - \rho) \theta\end{aligned}$$

```
16:   end for
17: end if
18: until convergence
```

Results

The following diagram shows the results after the DDPG algorithm is trained-



Episode: 10 Average Score: 0.75 Current Score: 0.84
 Episode: 20 Average Score: 0.94 Current Score: 1.50
 Episode: 30 Average Score: 3.05 Current Score: 11.66
 Episode: 40 Average Score: 5.54 Current Score: 13.41
 Episode: 50 Average Score: 7.56 Current Score: 18.70
 Episode: 60 Average Score: 10.33 Current Score: 31.77
 Episode: 70 Average Score: 12.77 Current Score: 25.41
 Episode: 80 Average Score: 14.23 Current Score: 24.11
 Episode: 90 Average Score: 15.58 Current Score: 30.80
 Episode: 100 Average Score: 16.99 Current Score: 29.24
 Episode: 110 Average Score: 20.12 Current Score: 32.69
 Episode: 120 Average Score: 23.16 Current Score: 31.80
 Episode: 130 Average Score: 25.61 Current Score: 30.23
 Episode: 140 Average Score: 27.37 Current Score: 30.45
 Episode: 150 Average Score: 28.95 Current Score: 31.97
 Episode: 160 Average Score: 29.70 Current Score: 33.94
 Episode: 167 Average Score: 30.05 Current Score: 32.61
 Environment solved in 67 episodes! Average Score: 30.05

The hyperparameters used in this model are given below-

Hyperparameter	Value
Number of Actions	4
Number of States	33
Number of Episodes	2000
Max time steps per episode	1000
Replay Buffer Size	1e6
Batch Size	1024
Soft update of target parameters	1e-3
Gamma	0.99
Actor learning rate	1e-3
Critic learning rate	1e-3
Leaky ReLU leakiness	0.01
D4PGActivation	Leaky ReLU
Update after every step	10

Improvements

- Other algorithms such as TRPO, PPO, A3C, D4PG could potentially lead to better results
- Using prioritized experience replay can improve the performance of the model
- Using a combination of on-policy and off-policy algorithm could potentially lead to better results. One such algorithm is Q-prop
- Using convoluted neural network architectures with complex activation functions could potentially give better results!